

Production Search Typeahead System

HLD Assignment Project Report

Student: Aditya Kumar Rai
SST Email: aditya.24bcs10178@sst.scaler.com
Stack: Java 21, Spring Boot, PostgreSQL, OpenSearch, Redis, Kafka

1. Architecture Diagram And Explanation

The system is designed as a production-style typeahead service where the API path is stateless, read-heavy prefix requests are cached, and search-count writes are durably buffered before batched updates.

```
Browser UI
-> Spring Boot API
  -> Suggestion Service
    -> Query Normalizer
    -> ConsistentHashRing
    -> Redis Prefix Cache Shards
    -> OpenSearch Suggestion Index
  -> Search Submission Service
    -> Kafka Search Events
  -> Batch Aggregator Worker
    -> PostgreSQL query-count upsert
    -> Trend bucket update
    -> OpenSearch weight refresh
    -> Redis prefix cache invalidation
  -> Metrics and Actuator
```

Component	Responsibility	Production Choice
Spring Boot API	REST APIs, validation, static UI, OpenAPI docs, metrics	Java 21 + Spring Boot
Query Store	Durable query text, historical counts, import metadata, trend buckets	PostgreSQL
Suggestion Index	Low-latency prefix retrieval and weighted ranking	OpenSearch completion suggester
Prefix Cache	Hot prefix result cache with TTL	Redis shards
Cache Router	Maps prefix keys to cache nodes	Java consistent hash ring with virtual nodes

Write Buffer	Durable search-event ingestion	Kafka
Batch Worker	Aggregates repeated searches and reduces DB writes	Spring scheduled/Kafka consumer worker

Suggestion Read Path

1. User types a prefix such as `iph`.
2. API normalizes the prefix and builds `suggest:v1:{version}:{rank}:{prefix}`.
3. Consistent hash ring maps the cache key to a cache node.
4. Cache hit returns top 10 suggestions immediately.
5. Cache miss queries the suggestion index, caches the result with TTL, and returns debug metadata.

Search Write Path

1. User submits a query through `POST /search`.
2. API validates and normalizes the query and publishes a search event.
3. Batch worker aggregates repeated queries every 2 seconds or 500 events.
4. Worker updates query counts, trend buckets, OpenSearch weights, and invalidates affected cache prefixes.

2. Dataset Source And Loading Instructions

Primary dataset: **AmazonQAC**, a real query-autocomplete dataset from Amazon Search logs.

- Dataset URL: <https://huggingface.co/datasets/amazon/AmazonQAC>
- License: CDLA-Permissive-2.0
- Fields used: `final_search_term`, `popularity`, `search_time`, `prefixes`
- Requirement coverage: train split is much larger than the required 100,000 queries.

The Java importer accepts AmazonQAC CSV or JSONL export files with these fields:

```
final_search_term,popularity,search_time,prefixes
```

Load at least 100,000 unique queries:

```
mvn -Dmaven.repo.local=.m2/repository spring-boot:run \
  -Dspring-boot.run.arguments="--typeahead.dataset.import-
  path=/path/to/amazonqac_100k.csv --typeahead.dataset.import-limit=100000"
```

The repository also includes `data/sample-amazonqac-smoke.csv` only for quick smoke testing. The official dataset source remains AmazonQAC.

3. API Documentation

Swagger/OpenAPI is available at `/docs` when the app is running.

API	Purpose	Notes
GET /suggest? q=iph&rank=count	Basic top 10 suggestions	Sorts by historical count only.
GET /suggest? q=iph&rank=hybrid	Trending-aware suggestions	Combines historical and recent activity.
POST /search	Dummy search submission	Returns Searched and queues count update.
GET /trending?window=1h	Trending searches	Supports 1h, 24h, 7d.
GET /cache/debug? prefix=iph	Cache routing debug	Shows selected node, hit/miss, key, TTL, ring settings.
GET /metrics	Performance evidence	Latency, p95, cache hit rate, DB writes, write reduction.

Example suggestion response:

```
{
  "query": "iph",
  "rank": "hybrid",
  "suggestions": [
    {
      "query": "iphone 15",
      "historical_count": 850000,
      "recent_count_1h": 312,
      "recent_count_24h": 1205,
      "score": 918240
    }
  ],
  "cache": {
    "hit": true,
    "node": "cache-node-2",
    "key": "suggest:v1:1:hybrid:iph"
  },
  "latency_ms": 3.7
}
```

4. Design Choices And Trade-offs

PostgreSQL Source Of Truth

PostgreSQL provides durable upserts, indexing, auditability, and easy inspection. It is not the hot read path; Redis and OpenSearch handle typeahead latency. At very large scale, the counter store could be moved to Cassandra, ScyllaDB, or DynamoDB.

OpenSearch Suggestion Index

OpenSearch completion suggesters provide prefix autocomplete with weighted ranking. The trade-off is eventual consistency because weights are refreshed after batch flushes.

Redis Cache With Consistent Hashing

Prefixes are read-heavy and highly repetitive, so caching top suggestions reduces latency. A Java consistent hash ring with 128 virtual nodes per cache node routes keys. Adding/removing a node remaps only a fraction of keys, unlike modulo hashing.

Trending Scoring

Basic mode uses only historical count. Enhanced mode uses:

```
score =  
  log10(historical_count + 1) * 100000  
  + log10(recent_count_1h + 1) * 60000  
  + log10(recent_count_24h + 1) * 25000
```

Recent activity is tracked in 5-minute buckets. Old spikes stop dominating because old buckets fall out of the selected window and hybrid cache TTL is short.

Kafka Batch Writes

Kafka is used as a durable write buffer. The batch worker aggregates repeated queries and performs fewer database writes. This trades immediate consistency for lower write pressure and better reliability than an in-memory-only buffer.

5. Performance Report

Smoke-profile benchmark command:

```
./scripts/benchmark.sh
```

Metric	Measured Value
Suggest requests	203
Cache hits	192
Cache misses	11
Cache hit rate	94.58%
Average suggest latency	0.16 ms
P95 suggest latency	0.34 ms
Search events received	540
DB write operations	7

Writes saved by batching	533
Write reduction	98.70%

These numbers were measured in the default smoke profile. The same benchmark script should be rerun after loading the 100,000-query AmazonQAC subset for final full-dataset evidence.

Demo Evidence

The UI demonstrates the required ranking difference. In historical mode, low cost flights wins by all-time count. After repeated searches, trending-aware mode lifts low latency search.

Search Typeahead

Production-style autocomplete with cache routing, trending, and batched writes.

Overall

Trending-aware

low

Search

low cost flights

190,000 score

low latency search

20,540 score

Trending

1h

1. low latency search

540 recent, 20,540 total

Cache Debug

Refresh

Rank

count

Hit

false

Node

cache-node-3

Key

suggest:v1:1:count:low

TTL

179s

Latency

0.8 ms

Metrics

Refresh

Hit Rate

93.7%

P95 Latency

0.4 ms

Search Events

540

Batch Flushes

7

Avg Latency

0.2 ms

Write Saved

98.7%

DB Writes

7

Invalidations

252

Search Typeahead

Production-style autocomplete with cache routing, trending, and batched writes.

Overall

Trending-aware

low

×

Search

low latency search

663,584 score

low cost flights

527,876 score

Trending

1h

1. low latency search

540 recent, 20,540 total

Cache Debug

Refresh

Rank

hybrid

Hit

false

Node

cache-node-2

Key

suggest:v1:1:hybrid:low

TTL

14s

Latency

0.7 ms

Metrics

Refresh

Hit Rate

94.1%

P95 Latency

0.3 ms

Search Events

540

Batch Flushes

7

Avg Latency

0.2 ms

Write Saved

98.7%

DB Writes

7

Invalidations

252

file:///Users/adityakumarrai/.openclaw/workspace/hld-search-suggestions-assignment/docs/project-report.html

6/7

Rubric Coverage

Component	Marks	Coverage
Basic Implementation	60	Dataset import, UI, suggestions API, search API, count updates, distributed cache with consistent hashing.
Trending Searches	20	Windowed buckets, hybrid score, /trending endpoint, stale-spike prevention, visible count-vs-hybrid demo.
Batch Writes	20	Kafka/event-buffer design, aggregation, write-reduction metrics, failure trade-off explanation.